

1st Winter primitive is a streaming chunked AEAD with CMT-4.

- AEAD is AEAD(K, N, A, M) where K=key, N=nonce, A=associated data, M=message
- CMT-1 is AEAD that commits to K only
- CMT-4 is AEAD that commits to all 4 (K, N, A, M)

GCM-based, with <https://eprint.iacr.org/2022/268.pdf> and <https://eprint.iacr.org/2020/1456.pdf>  
AEAD (GCM) → CMT-1 (key-commitment tag) → CMT-4 (commitment to K, N, A, M)

GCM → CMT-1:

**Algorithm 2:**  $\text{CommitKey}_{IV} \Pi^{\text{Enc}}(K, N', N, A, M)$

**Input:**  $K \in \{0, 1\}^{k_0}, N' \in \{0, 1\}^{v'}, N \in \mathcal{N}, A \in \mathcal{A}, M \in \mathcal{M}$

**Output:**  $C \in \mathcal{C}, K_{\text{com}} \in \{0, 1\}^c$

1  $K_{\text{enc}} \leftarrow F_{\text{enc}}(K, N')$

2  $K_{\text{com}} \leftarrow F_{\text{com}}(K, N')$

3  $C \leftarrow \text{Enc}(K_{\text{enc}}, N, A, M)$

4 **return**  $(C, K_{\text{com}})$

CMT-1 → CMT-4 (MB/VTH original scheme):

|   |   |
|---|---|
| $\overline{\text{SE}}.\text{Enc}(K, N, A, M)$             | $\overline{\text{SE}}.\text{Dec}(K, N, A, C^* \  T')$     |
| $L \leftarrow H(K, (N, A))$                               | $L \leftarrow H(K, (N, A))$                               |
| $C \leftarrow \text{SE}.\text{Enc}(L, N, \varepsilon, M)$ | $M \leftarrow \text{SE}.\text{Dec}(L, N, \varepsilon, C)$ |
| Return $C$  | Return $M$  |

Using a combined version for GCM → CMT-4 (with some tweaks):

Inputs: User Key [UK:\*], Plaintext [PTX:\*], optional Associated Data [AD:\*], Chunk Length CL=constant

- HASH is SHA512
- HKDF is HKDF<HASH>
- GCM is GCM with 256-bit key and 128-bit tag
- [CN:8] is little-endian 8-byte UInt64

Encrypt per-Message:

1. [MID:40] ← Random Byte Generator RBG:40 (MID is Message ID)

2. [ADH:64] ← HASH([AD:\*])

3. [KDK:64] ← HKDF.Extract(ikm: [UK:\*], salt: ["KDK"][MID:40][ADH:64]) // salt not to exceed 128 bytes

4. [MCT:32] ← HKDF.Expand(prk: [KDK:64], len: 32, info: ["MCT"][MID:40][ADH:64]) // info not to exceed 109 bytes

5. [MHR:72] ← [MID:40][MCT:32] // MCT is Msg Commitment Tag

6. Output Message Header [MHR:72]

7. Keep in internal per-message state: [KDK:64], [MHR:72]

Encrypt per-Chunk:

1. [CEK:32][CIV:12] ← HKDF.Expand(prk: [KDK:64], len: 44, info: ["CEI"][MHR:72][CN:8])

2. GCM.Enc(key:[CEK:32], iv:[CIV:12], ptx:[PTX:CL], ctx: out [CTX:CL], ad:[empty], tag: out [GT:16])

3. Return [CTX:CL][GT:16]

1st **Winter** primitive -- streaming chunked AEAD -- output:

| Header | Chunk 1 | Chunk ... | Chunk L (Last) |
|--------|---------|-----------|----------------|
|--------|---------|-----------|----------------|

- Header is [MID:40][MCT:32]
- Full Chunk is 128 KiB ( $1024 \times 128 = 131,072$  bytes)
  - [131,056 bytes of GCM ciphertext][16 bytes of GCM tag]
- Last Chunk is strictly  $< 131,072$  bytes (Chunk 1 can be Last Chunk)
- Encrypted message can have up to  $2^{64}$  chunks
  - [Maximum](#) message size is 2 YiB (Yobibytes). 1 Yobibyte =  $2^{80}$  bytes.
- $2^{128}$  Messages can be created with MID collision probability  $1/2^{64}$ .
- Per-Chunk GCM-tags provide [~115 bits](#) of security (NIST requires no fewer than 112 bits)..

**Winter** provides primitives for building Cloud Scale Encryption Systems (described by [Campagna/Gueron 2019](#)):

- Any (very large) number of users
- Any (very large) number of messages
- Any (very large) number of keys
- Any (very large) number of encryptions per key
- Any (very large) maximum message size

**Winter – Cloud Scale Cryptography Toolkit** discussion:

1st **Winter** primitive -- streaming chunked AEAD

1. Design motivation (why bother creating it?):
  - [Adam Langley](#) blog (Senior Staff Engineer @Google)
  - “How do I encrypt lots of records with a per-user key?” – (every) Anonymous developer
  - This symmetric-cryptography question is the original motivation of **ALL OF CRYPTOGRAPHY**.
  - You would have thought that we would have figured it out by now, but we totally haven't.
  - This was true in 2015 and is still true in 2022.
  - Cloud Scale cryptographic systems changed the definition of “lots”.
2. Design goals (what are we aiming to achieve?):
  - Streaming chunked [AEAD](#) with commitment to all 4 inputs (Key, Nonce, Associated Data, Message)
  - Cloud Scale Cryptography – safe security bounds (beyond AES birthday bounds) under very large:
    - plaintext sizes, # of users/keys/operations per key/forgery attempts
  - Utmost safety & usability (“[Security at the expense of usability comes at the expense of security.](#)”)
  - Regulatory Compliance Ready (targeting NIST as well as international compliance requirements)
    - “The best cryptography is useless if it cannot be used.”
  - Implementation interoperability across modern languages/frameworks
    - .NET, Java, Golang, Python, Swift, JS, Rust, C/C++, PHP, etc.
    - (ex. sp800\_108 KDFs are rarely available, but HKDF is ubiquitous)
  - Fast speed (throughput), targeting at least 3 GiB/second on 1 average CPU core
  - Encryption and Decryption speed should scale with additional CPUs/memory
  - Single ciphertext byte-format (regardless of CPUs/memory used, or whether streaming or not)
    - addresses usability requirement
  - Parallelizable chunk encryption, random chunk access (decryption)
  - Every byte of output should be (1) indistinguishable from random; (2) authenticated
  - Minimal plaintext-to-ciphertext size expansion
3. Cryptographic primitives used:
  - **AES-GCM** ([NIST SP 800-38D](#)) with 256-bit key and 128-bit tag
  - **SHA-512** ([RFC 6234](#), [FIPS 180-4](#))
  - **HKDF** ([RFC 6234](#), [NIST SP 800-56C](#), [NIST ACVP](#)), which uses HMAC ([RFC 5869](#), [FIPS 198-1](#))

---

## Winter – Cloud Scale Cryptography Toolkit – FAQ

1. Why use it?
  - a. **Extreme performance.** Early streaming-AEAD prototypes run at ~**6.5 GiB/second** on a 4-CPU Intel.
  - b. Usable cryptography. No limits to worry about, no cryptographic-wearout of anything, no need to rekey.
  - c. No choices/decisions to make. Only 1 protocol per primitive, no interoperability issues.
  - d. Designed for regulatory compliance (safe to use in your Business/Government/Industry).
2. Who is designing it?
  - a. Stan Drapkin is responsible for the design (and any mistakes), but design ideas were shaped by discussions with cryptography experts. Your feedback is appreciated.
3. Is it Production-ready?
  - a. No (not yet), but at some point it will be. The 1st stage is to get as much feedback as possible.
4. Is implementation ready?
  - a. If early feedback does not identify big issues, an MVP/reference implementation will be shared on [Github](#).
5. Why is Chunk size 128 KiB?
  - a. 128 KiB chunks achieve the best performance (throughput) in our tests, while balancing other factors.
6. How does **Winter** streaming AEAD compare with [Google Tink](#)? (not a comprehensive comparison)
  - a. Tink provides crypto-agility. Winter has no cryptographic agility by design (only 1 protocol per primitive).
  - b. Tink's AEADs (plural) use a same-key GCM for all chunks of a message (changing Nonce only).
  - c. Tink's limits are inferior to Winter's AEAD limits.
  - d. Tink's performance is inferior to Winter's AEAD performance.
  - e. Tink has no official .NET support.
  - f. Tink AEADs are not CMT-1 or CMT-4.
7. How does **Winter** streaming AEAD compare with [AWS Encryption SDK](#)? (not a comprehensive comparison)
  - a. AESDK provides crypto-agility. Winter has no cryptographic agility by design (only 1 protocol per primitive).
  - b. AESDK AEADs use a same-key GCM for all chunks of a message (changing Nonce only).
  - c. AESDK limits are inferior to Winter's AEAD limits.
  - d. AESDK performance is inferior to Winter's AEAD performance.
  - e. AESDK does not support streaming for .NET.
  - f. AESDK AEADs are CMT-1 but may not be CMT-4.